

```

import pandas as pd, numpy as np, matplotlib.pyplot as plt, warnings, os, joblib
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.metrics import (accuracy_score, precision_score, recall_score, f1_score,
                             roc_auc_score, classification_report, confusion_matrix, roc_curve, auc)
warnings.filterwarnings('ignore')
plt.rcParams['figure.figsize'] = (8,5)

```

```

data=pd.read_csv('/content/Lung_Cancer.csv')
df=pd.DataFrame(data)
print('Loaded dataset with shape:',df.shape)
df.head()

```

Loaded dataset with shape: (890000, 17)

	id	age	gender	country	diagnosis_date	cancer_stage	family_history	smoking_status	bmi	cholesterol_level	hypertension
0	1	64	Male	Sweden	05-04-2016	Stage I		Yes	Passive Smoker	29.4	199
1	2	50	Female	Netherlands	20-04-2023	Stage III		Yes	Passive Smoker	41.2	280
2	3	65	Female	Hungary	05-04-2023	Stage III		Yes	Former Smoker	44.0	268
3	4	51	Female	Belgium	05-02-2016	Stage I		No	Passive Smoker	43.0	241
4	5	37	Male	Luxembourg	29-11-2023	Stage I		No	Passive Smoker	19.7	178

```

print('Dtypes:')
df.dtypes

```

Dtypes:

	0
id	int64
age	int64
gender	object
country	object
diagnosis_date	object
cancer_stage	object
family_history	object
smoking_status	object
bmi	float64
cholesterol_level	int64
hypertension	int64
asthma	int64
cirrhosis	int64
other_cancer	int64
treatment_type	object
end_treatment_date	object
survived	int64

dtype: object

```

print('\nMissing values:')
print(df.isnull().sum().sort_values(ascending=False).head(30))

```

Missing values:

id	0
age	0
gender	0
country	0
diagnosis_date	0
cancer_stage	0

```

family_history      0
smoking_status     0
bmi                0
cholesterol_level  0
hypertension        0
asthma              0
cirrhosis           0
other_cancer        0
treatment_type      0
end_treatment_date 0
survived            0
dtype: int64

```

```
display(df.info())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 890000 entries, 0 to 889999
Data columns (total 17 columns):
 #   Column          Non-Null Count  Dtype  
--- 
 0   id              890000 non-null   int64  
 1   age             890000 non-null   int64  
 2   gender           890000 non-null   object  
 3   country          890000 non-null   object  
 4   diagnosis_date  890000 non-null   object  
 5   cancer_stage    890000 non-null   object  
 6   family_history   890000 non-null   object  
 7   smoking_status   890000 non-null   object  
 8   bmi              890000 non-null   float64 
 9   cholesterol_level 890000 non-null   int64  
 10  hypertension     890000 non-null   int64  
 11  asthma            890000 non-null   int64  
 12  cirrhosis         890000 non-null   int64  
 13  other_cancer      890000 non-null   int64  
 14  treatment_type   890000 non-null   object  
 15  end_treatment_date 890000 non-null   object  
 16  survived          890000 non-null   int64  
dtypes: float64(1), int64(8), object(8)
memory usage: 115.4+ MB
None

```

```
display(df.describe(include='all').T)
```

	count	unique	top	freq	mean	std	min	25%	50%	75%	89
id	890000.0	NaN		NaN	445000.5	256921.014128	1.0	222500.75	445000.5	667500.25	89
age	890000.0	NaN		NaN	55.007008	9.994485	4.0	48.0	55.0	62.0	
gender	890000	2	Male	445134		NaN	NaN	NaN	NaN	NaN	
country	890000	27	Malta	33367		NaN	NaN	NaN	NaN	NaN	
diagnosis_date	890000	3651	15-05-2024	306		NaN	NaN	NaN	NaN	NaN	
cancer_stage	890000	4	Stage III	222594		NaN	NaN	NaN	NaN	NaN	
family_history	890000	2	No	445181		NaN	NaN	NaN	NaN	NaN	
smoking_status	890000	4	Passive Smoker	223170		NaN	NaN	NaN	NaN	NaN	
bmi	890000.0	NaN		NaN	30.494172	8.368539	16.0	23.3	30.5	37.7	
cholesterol_level	890000.0	NaN		NaN	233.633916	43.432278	150.0	196.0	242.0	271.0	
hypertension	890000.0	NaN		NaN	0.750024	0.432999	0.0	1.0	1.0	1.0	
asthma	890000.0	NaN		NaN	0.46974	0.499084	0.0	0.0	0.0	0.0	
cirrhosis	890000.0	NaN		NaN	0.225956	0.418211	0.0	0.0	0.0	0.0	
other_cancer	890000.0	NaN		NaN	0.088157	0.283524	0.0	0.0	0.0	0.0	
treatment_type	890000	4	Chemotherapy	223262		NaN	NaN	NaN	NaN	NaN	
end_treatment_date	890000	4194	09-12-2023	294		NaN	NaN	NaN	NaN	NaN	
...

```

dups = df.duplicated().sum()
print('\nNumber of duplicate rows:', dups)

```

```
Number of duplicate rows: 0
```

```
df['survived'].value_counts()
```

```
count
```

```
survived
```

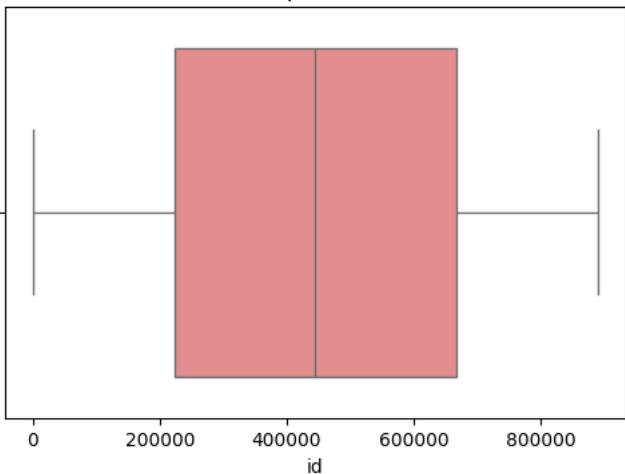
0	693996
1	196004

```
dtype: int64
```

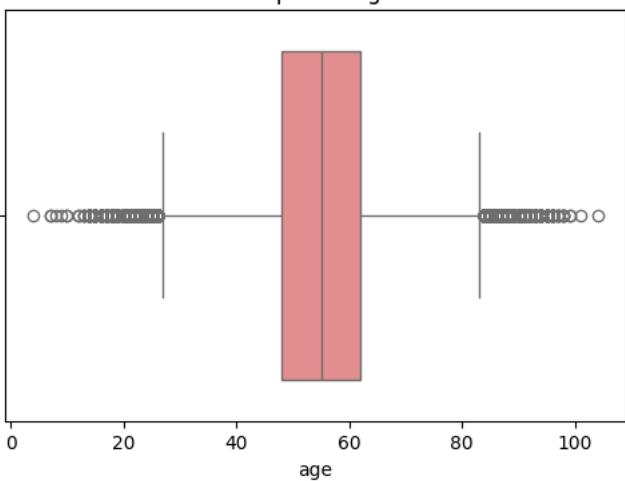
```
num_cols = df.select_dtypes(include=['int64', 'float64']).columns.tolist()
```

```
import seaborn as sns
top_num = num_cols[:6] if len(num_cols) > 6 else num_cols
for col in top_num:
    plt.figure(figsize=(6,4))
    sns.boxplot(x=df[col], color='lightcoral')
    plt.title(f'Boxplot of {col}')
    plt.show
```

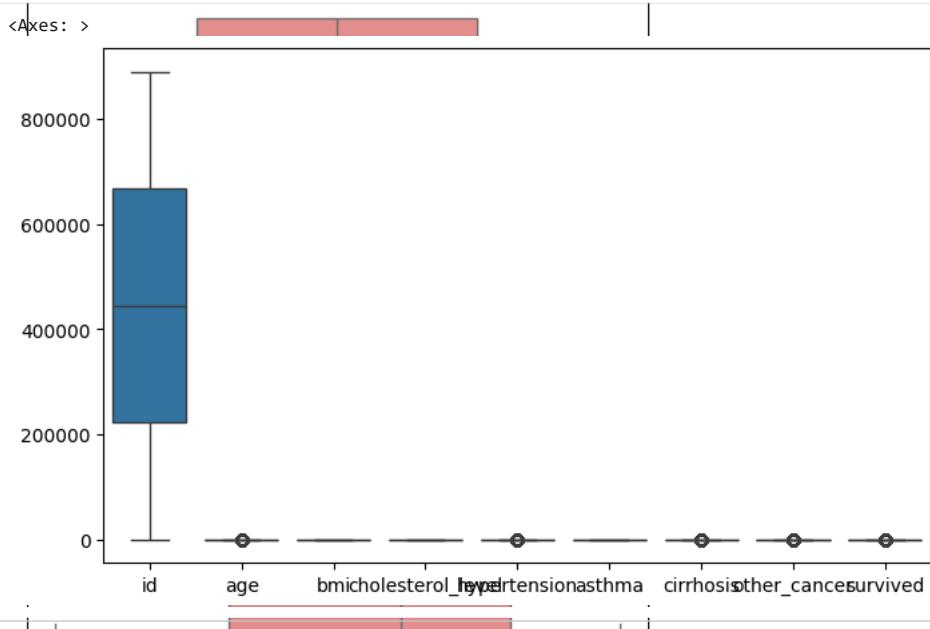

Boxplot of id



Boxplot of age



```
sns.boxplot(data=df)
```



```
def handling_outliers(df, exclude_cols=None):
    if exclude_cols is None:
        exclude_cols = []

    for col in df.select_dtypes(include='number').columns:
        if col in exclude_cols:
            continue # skip target or excluded columns

        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - (1.5 * IQR)
        upper_bound = Q3 + (1.5 * IQR)
```

```

# Capping outliers
df[col] = df[col].apply(lambda x: lower_bound if x < lower_bound
                        else upper_bound if x > upper_bound else x)
return df

```

```

df1 = handling_outliers(df, exclude_cols=['survived'])
df1

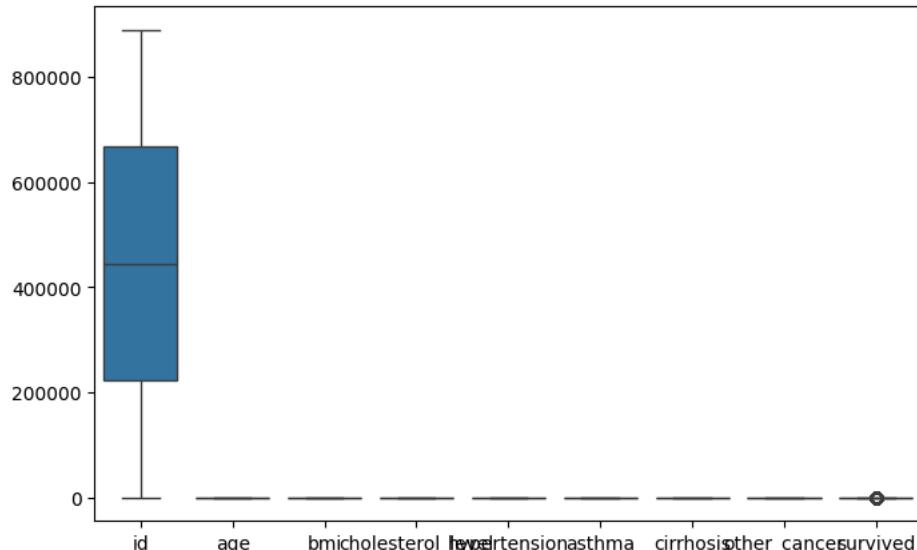
```

	id	age	gender	country	diagnosis_date	cancer_stage	family_history	smoking_status	bmi	cholesterol_lev
0	1	64.0	Male	Sweden	05-04-2016	Stage I	Yes	Passive Smoker	29.4	1
1	2	50.0	Female	Netherlands	20-04-2023	Stage III	Yes	Passive Smoker	41.2	2
2	3	65.0	Female	Hungary	05-04-2023	Stage III	Yes	Former Smoker	44.0	2
3	4	51.0	Female	Belgium	05-02-2016	Stage I	No	Passive Smoker	43.0	2
4	5	37.0	Male	Luxembourg	29-11-2023	Stage I	No	Passive Smoker	19.7	1
hypertension										
...
889995	889996	40.0	Male	Malta	01-07-2022	Stage IV	No	Passive Smoker	44.8	2
889996	889997	62.0	Female	Cyprus	27-09-2015	Stage II	Yes	Former Smoker	21.6	2
889997	889998	48.0	Female	Estonia	27-03-2016	Stage III	Yes	Never Smoked	38.6	2
889998	889999	67.0	Female	Slovakia	22-12-2015	Stage IV	Yes	Former Smoker	18.6	1
889999	890000	55.0	Female	Malta	26-07-2021	Stage II	Yes	Current Smoker	42.8	2

890000 rows × 17 columns

```
sns.boxplot(data=df1)
```

```
<Axes: >
```



```
df1['survived'].value_counts()
```

```
count
```

```
survived
```

survived	count
0	693996
1	196004

```
dtype: int64
```

```

from sklearn.preprocessing import LabelEncoder
df_encoded = df.copy()
categorical_cols = df_encoded.select_dtypes(include=['object', 'category']).columns.tolist()
le = LabelEncoder()
for col in categorical_cols:
    df_encoded[col] = le.fit_transform(df_encoded[col].astype(str))

print("Label encoding complete!")

```

```
print("Encoded DataFrame shape:", df_encoded.shape)
df_encoded.head()
```

Label encoding complete!

Encoded DataFrame shape: (890000, 17)

	id	age	gender	country	diagnosis_date	cancer_stage	family_history	smoking_status	bmi	cholesterol_level	hypertension	asthma	cirrhosis	other_cancer	treatment_type	end_treatment_date	survived
0	1	64.0	1	26	510	0	1	3	29.4	199	0	0	0	0	0	0	1
1	2	50.0	0	19	2317	2	1	3	41.2	280	1	0	0	0	0	0	0
2	3	65.0	0	12	517	2	1	1	44.0	268	1	0	0	0	0	0	0
3	4	51.0	0	1	490	0	0	3	43.0	241	1	0	0	0	0	0	0
4	5	37.0	1	17	3461	0	0	3	19.7	178	1	0	0	0	0	0	0

```
from imblearn.over_sampling import SMOTE

X = df_encoded.drop(columns=['survived'])
y = df_encoded['survived']

smote = SMOTE(random_state=42)
X_balanced, y_balanced = smote.fit_resample(X, y)

print(y_balanced.value_counts())

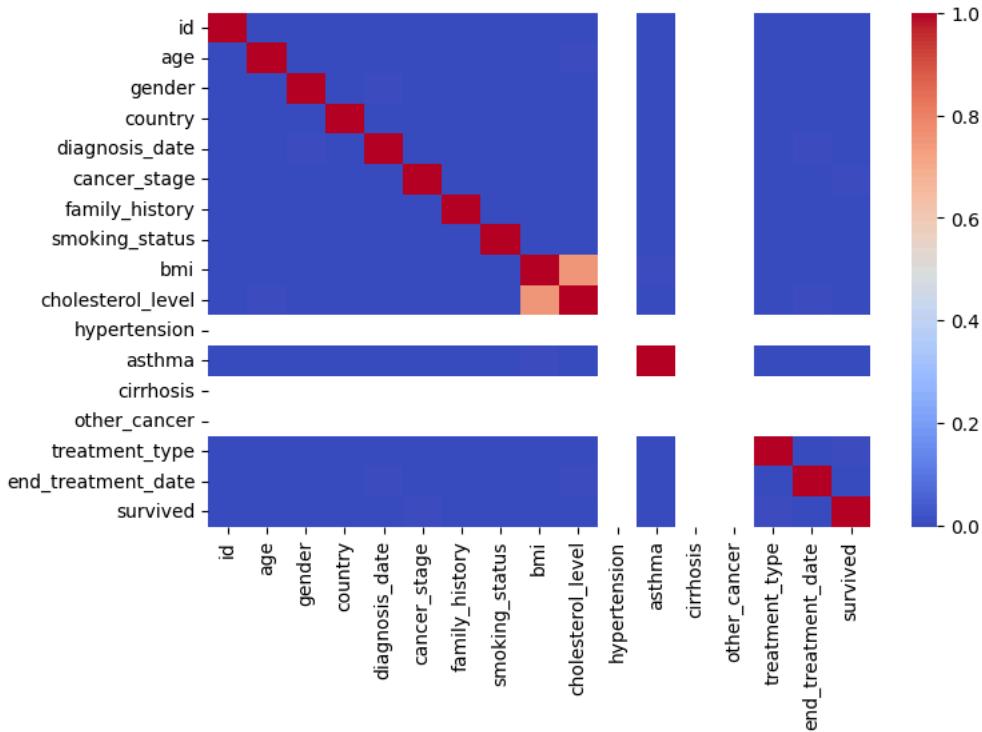
survived
0    693996
1    693996
Name: count, dtype: int64
```

```
c=df_encoded.corr(numeric_only=True)
c
```

	id	age	gender	country	diagnosis_date	cancer_stage	family_history	smoking_status	bmi	cholesterol_level	hypertension	asthma	cirrhosis	other_cancer	treatment_type	end_treatment_date	survived	
id	1.000000	0.000135	0.000079	0.000671	-0.001672	-0.001202	-0.000245	0.000236	-0.	0.0001507	0.001814	0.000224	-0.000483	-0.001084	-0.000504	0.000086	0.001169	0.
age	0.000135	1.000000	-0.000954	-0.000316	0.000856	0.000912	-0.001044	-0.000013	0.	0.000096	-0.000552	0.001142	-0.000521	-0.001362	-0.001350	-0.000808	-0.000140	0.
gender	0.000079	-0.000954	1.000000	0.001468	0.001933	-0.000161	0.000762	-0.000595	0.	0.000097	0.000181	0.000224	-0.000483	-0.001084	-0.000504	0.000086	0.001169	0.
country	0.000671	-0.000316	0.001468	1.000000	0.001281	0.000000	-0.000665	-0.000429	0.	0.000098	0.000182	0.000225	-0.000484	-0.001085	-0.000505	0.000087	0.001170	0.
diagnosis_date	-0.001672	0.000856	0.001933	0.001281	1.000000	-0.000429	-0.000518	0.000345	0.	0.000099	0.000183	0.000226	-0.000485	-0.001086	-0.000506	0.000088	0.001171	0.
cancer_stage	-0.001202	0.000912	-0.000161	-0.000665	-0.000429	1.000000	0.000345	0.000000	0.	0.000099	0.000183	0.000226	-0.000485	-0.001086	-0.000506	0.000088	0.001171	0.
family_history	-0.000245	-0.001044	0.000762	-0.000624	-0.000518	0.000345	1.000000	-0.000000	0.	0.000099	0.000183	0.000226	-0.000485	-0.001086	-0.000506	0.000088	0.001171	0.
smoking_status	0.000236	-0.000013	-0.000595	0.001484	0.000490	0.001160	-0.001261	1.000000	0.	0.000099	0.000183	0.000226	-0.000485	-0.001086	-0.000506	0.000088	0.001171	0.
bmi	-0.001144	0.000955	0.000504	-0.000766	-0.000474	-0.000446	0.000411	0.000901	1.	0.000099	0.000183	0.000226	-0.000485	-0.001086	-0.000506	0.000088	0.001171	0.
cholesterol_level	-0.001507	0.001814	0.000224	-0.000483	-0.001084	-0.000504	0.000086	0.001169	0.	0.000099	0.000183	0.000226	-0.000485	-0.001086	-0.000506	0.000088	0.001171	0.
hypertension	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
asthma	0.000096	-0.000552	0.001142	-0.000521	-0.001362	-0.001350	-0.000808	-0.000140	0.	0.000099	0.000183	0.000226	-0.000485	-0.001086	-0.000506	0.000088	0.001171	0.
cirrhosis	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
other_cancer	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
treatment_type	0.000369	-0.001032	0.000040	0.001365	0.000915	-0.002192	0.000310	-0.000910	0.	0.000099	0.000183	0.000226	-0.000485	-0.001086	-0.000506	0.000088	0.001171	0.
end_treatment_date	0.000207	-0.000936	-0.000081	0.000043	0.001885	0.000992	0.001249	0.001188	-0.	0.000099	0.000183	0.000226	-0.000485	-0.001086	-0.000506	0.000088	0.001171	0.
survived	0.000517	0.001271	0.000762	0.000066	-0.002406	0.002519	0.001322	0.000087	0.	0.000099	0.000183	0.000226	-0.000485	-0.001086	-0.000506	0.000088	0.001171	0.

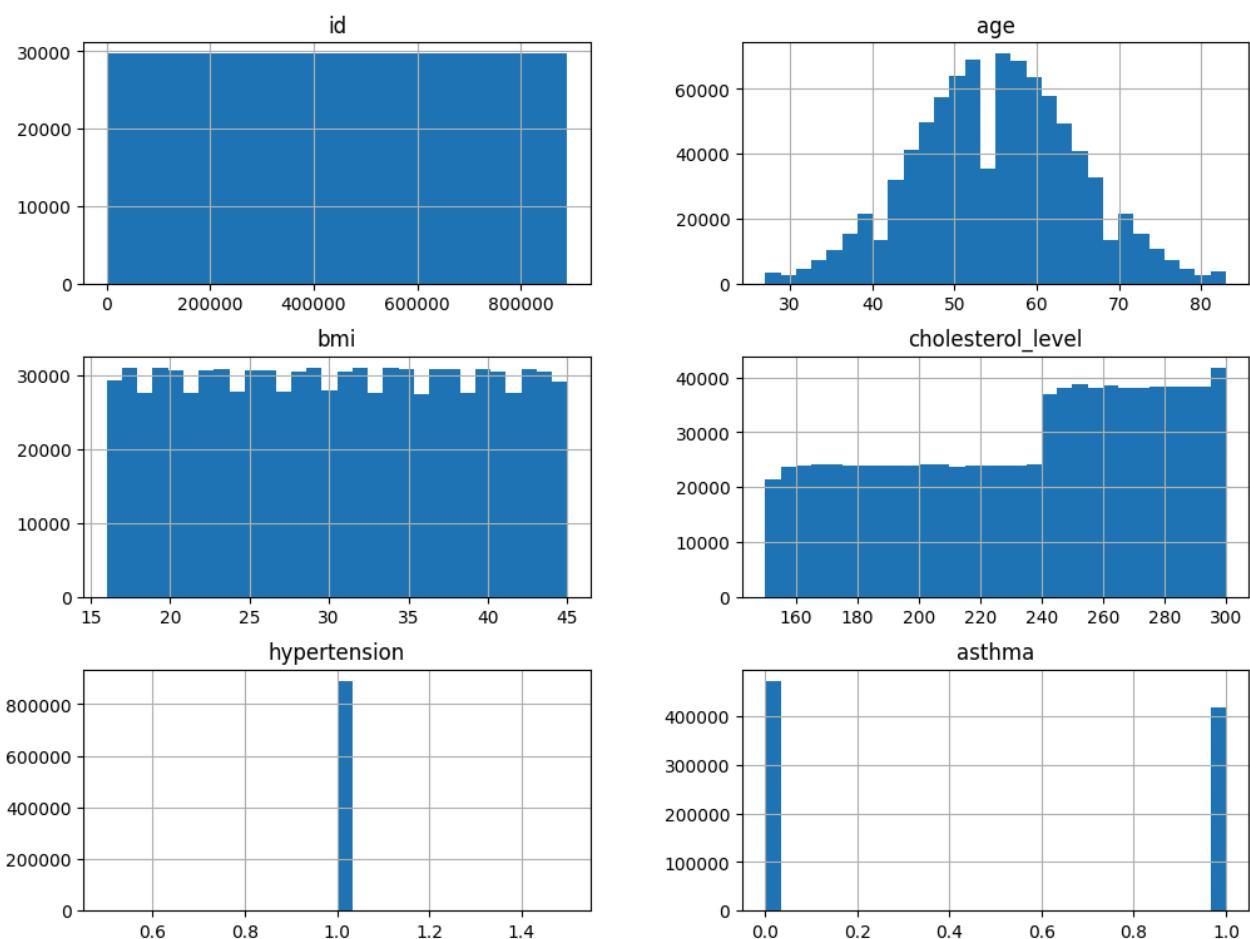
```
sns.heatmap(c,cmap='coolwarm')
```

```
<Axes: >
```



```
if len(num_cols)>0:  
    cols = num_cols[:6]  
    df1[cols].hist(bins=30, layout=(3,2), figsize=(12,9)); plt.suptitle('Numeric histograms'); plt.show()
```

Numeric histograms



```
for c in num_cols[:4]:
```

```
plt.figure(); sns.boxplot(x=df_encoded[c]); plt.title('Boxplot ' + c); plt.show()
```

```
X.columns
```

```
Index(['id', 'age', 'gender', 'country', 'diagnosis_date', 'cancer_stage',
       'family_history', 'smoking_status', 'bmi', 'cholesterol_level',
       'hypertension', 'asthma', 'cirrhosis', 'other_cancer', 'treatment_type',
       'end_treatment_date'],
      dtype='object')
```

```
y
```

```
survived
0      0
1      1
2      0
3      0
4      0
...
889995    0
889996    0
889997    1
889998    0
889999    0
890000 rows × 1 columns
```

```
dtype: int64
```

```
df_encoded.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 890000 entries, 0 to 889999
Data columns (total 17 columns):
 #   Column            Non-Null Count  Dtype  
 ---  -- 
 0   id                890000 non-null   int64  
 1   age               890000 non-null   float64 
 2   gender             890000 non-null   int64  
 3   country            890000 non-null   int64  
 4   diagnosis_date    890000 non-null   int64  
 5   cancer_stage      890000 non-null   int64  
 6   family_history     890000 non-null   int64  
 7   smoking_status     890000 non-null   int64  
 8   bmi                890000 non-null   float64 
 9   cholesterol_level 890000 non-null   int64  
 10  hypertension        890000 non-null   float64 
 11  asthma              890000 non-null   int64  
 12  cirrhosis           890000 non-null   float64 
 13  other_cancer        890000 non-null   float64 
 14  treatment_type      890000 non-null   int64  
 15  end_treatment_date 890000 non-null   int64  
 16  survived            890000 non-null   int64  
dtypes: float64(5), int64(12)
memory usage: 115.4 MB
```

```

from sklearn.feature_selection import SelectKBest,f_classif
selector = SelectKBest(score_func=f_classif, k=13)

X_new = selector.fit_transform(X, y)
selected_features = X.columns[selector.get_support()]

print("Top Selected Features:")
print(selected_features)
Index(['id', 'age', 'gender', 'country', 'diagnosis_date', 'cancer_stage',
       'family_history', 'smoking_status', 'bmi', 'cholesterol_level',
       'asthma', 'treatment_type', 'end_treatment_date'],
      dtype='object')

print(y.value_counts())
survived
0    693996
1    196004
Name: count, dtype: int64

x_train,x_test,y_train, y_test =train_test_split(X,y,test_size=0.2,random_state=42)

ss=StandardScaler()
x_train_scaled=ss.fit_transform(x_train)
x_test_scaled=ss.transform(x_test)

models = {
    'LogisticRegression': LogisticRegression(),
    'RandomForest': RandomForestClassifier(),
    'GradientBoosting': GradientBoostingClassifier(random_state=42)
}

print(y_train.value_counts())
survived
0    555357
1    156643
Name: count, dtype: int64

results = {}
for name, model in models.items():
    print('\nTraining', name)
    model.fit(x_train_scaled, y_train)
    preds = model.predict(x_test_scaled)
    acc = accuracy_score(y_test, preds)
    print(f'Accuracy on test set: {acc:.4f}')
    print('Classification report:')
    print(classification_report(y_test, preds))
    results[name] = {'model': model, 'accuracy': acc}
res_df = pd.DataFrame({k: {accuracy: v['accuracy']} for k,v in results.items()}).T.sort_values('accuracy', ascending=False)
print('\nModel comparison:')
display(res_df)

```

Training LogisticRegression
Accuracy on test set: 0.7789